

# *High Throughput Workflow Tools and Strategies*

## Part 3: HyperShell

Geoffrey Lentner

*Lead Research Data Scientist*

*Rosen Center for Advanced Computing*



**PURDUE**  
UNIVERSITY®

Rosen Center for  
Advanced Computing

# 1

## Preface

About me, background, chapter goals

# *Rosen Center for Advanced Computing*

Research computing and data services at Purdue University



- *RCAC* for short, Part of *Purdue IT*
- Operate *Top500* supercomputers, including PB-scale storage systems.
- Scientific Applications
- Envision Center
- Research Software Engineering
- Visit [rcac.purdue.edu](http://rcac.purdue.edu)



# About Me

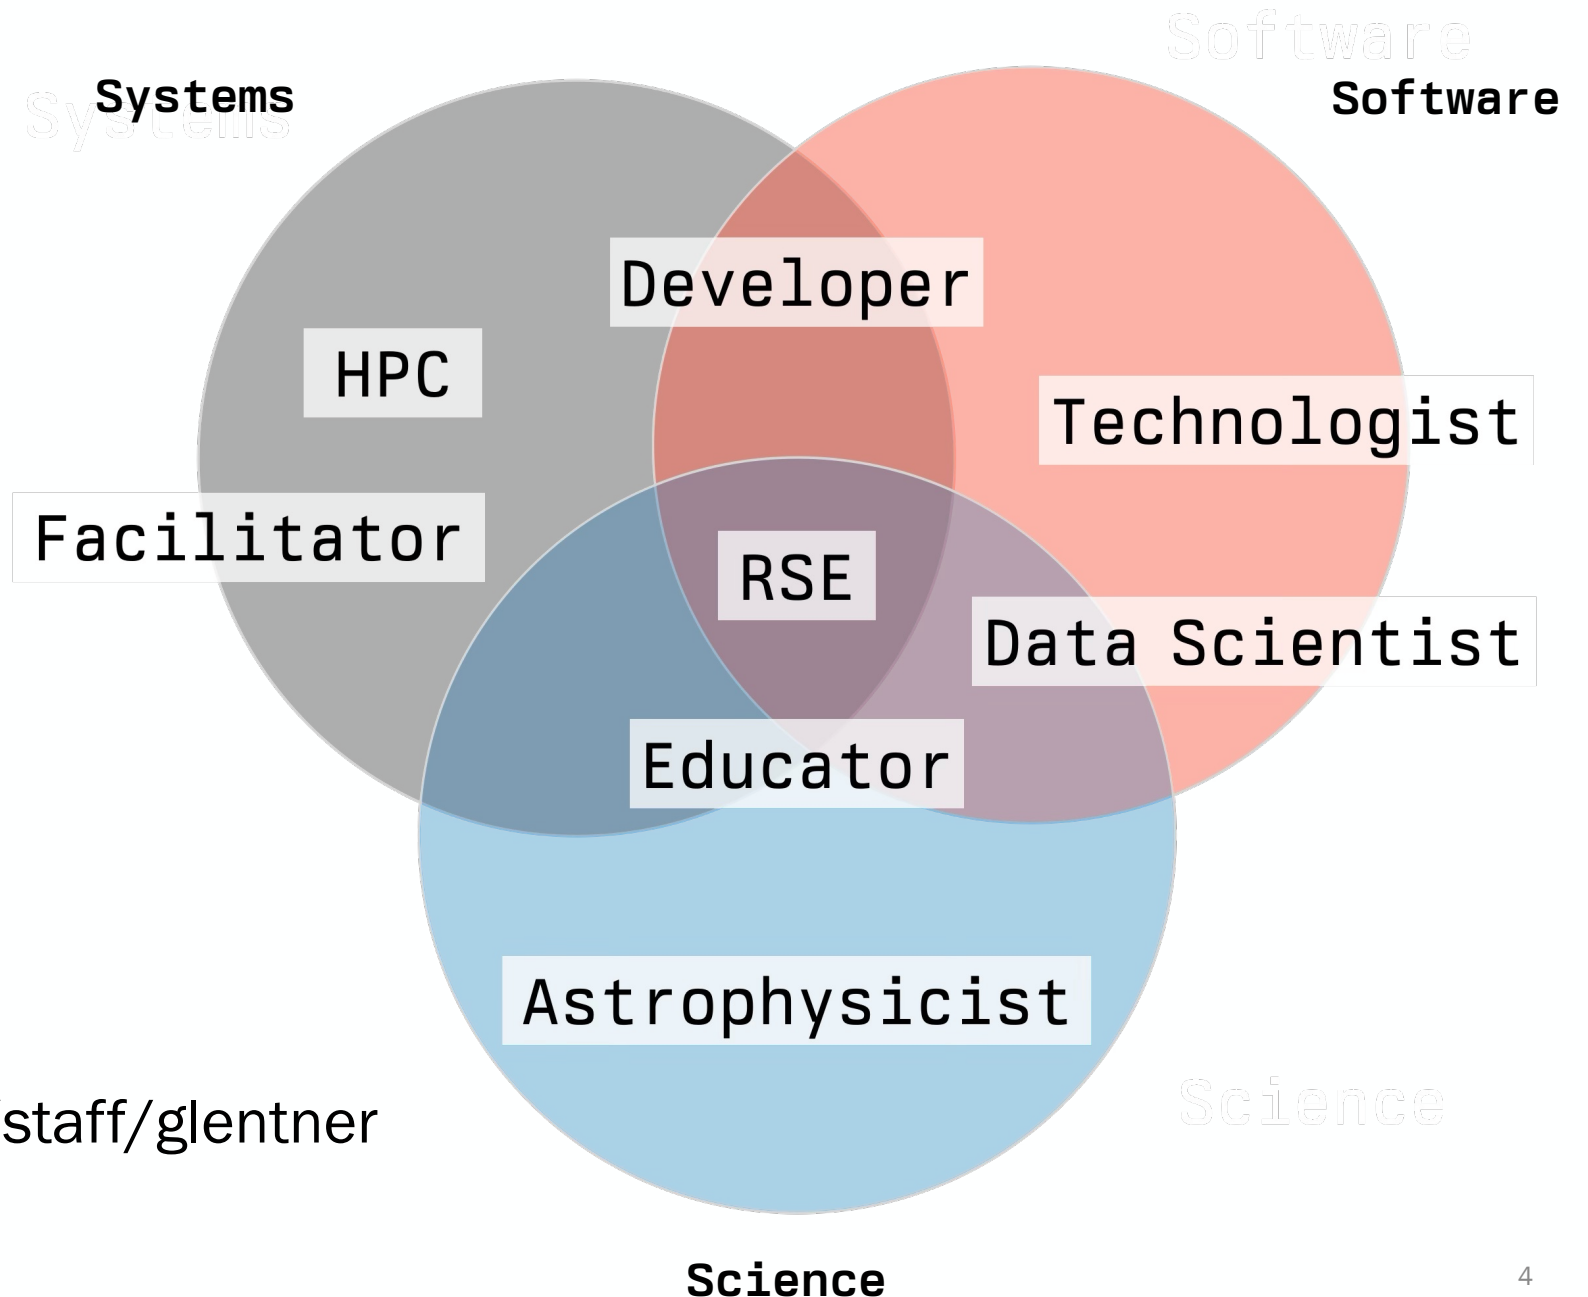
## Background and Roles



## Geoffrey Lentner

Email: [glentner@purdue.edu](mailto:glentner@purdue.edu)

Page: [rcac.purdue.edu/about/staff/glentner](http://rcac.purdue.edu/about/staff/glentner)



# Goals

What I hope to accomplish

- Speak on high-throughput computing issues.
- HyperShell features relative to other tools.
- Demonstrate various scenarios.

# 2

## Overview

Motivation; TLDR; project overview.

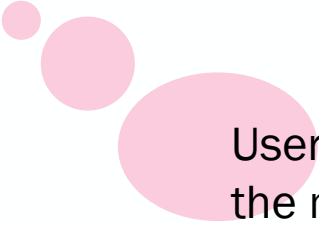
*Why use a workflow tool at all?*

*Why not use the main scheduler?*

# *Practical Limits*

Site administrators do not want users ...

- submitting *millions of jobs*,
- filling up the *database*,
- impacting site-wide *throughput*,
- polluting the *queue*.



Users don't want to be rate limited by the main schedulers bandwidth,

Or lose their job history over time.

...



**User:** *I want to run my job a million times...*

**Me:** *Um... About that...*

# HyperShell<sup>2.7.0</sup>

## The ultimate workflow automation tool

An elegant, cross-platform *high-throughput computing* utility for processing shell commands over a distributed, asynchronous queue.

A highly scalable workflow automation tool for many-task scenarios.



Windows




Mac



Linux



python



*But why a new tool?*

***TLDR;***

Very fancy shell processing ...



# ***TLDR;***

Combine the best features from tools in this genre

- Scale farther with task *aggregation*,
- Manage tasks with persistent *database*,
- Elastic (client-server) with *scale-to-zero*,
- *Cross-platform* (Linux server, Windows clients),
- Better *ergonomics* for researchers,
- Embed as a *library* in your project.

# ***TLDR;***

## Online resources

- *Code* at [github.com/hypershell/hypershell](https://github.com/hypershell/hypershell),
- *Docs* at [hypershell.readthedocs.io](https://hypershell.readthedocs.io),
- *Web* at [hypershell.org](https://hypershell.org), (under construction)
- *Discord* at [discord.gg/wmv5gyUfkN](https://discord.gg/wmv5gyUfkN).

# Getting Started

Installing as an unprivileged user (any platform)

See [installation](#) guide for details.

Optional dependency on *PostgreSQL* – can use *Anaconda* if necessary.

Install HyperShell on any platform

```
$ pip install hypershell  
$ pip install 'hypershell[postgres]'  
  
$ uv tool install 'hypershell[postgres]' --python 3.13
```



Install “uv” from <https://docs.astral.sh/uv/getting-started/installation/>

# *Getting Started*

## Installing on MacOS

Local installation on *MacOS* with *Homebrew*

Install HyperShell on macOS

```
$ brew tap hypershell/tap  
$ brew install hypershell
```



# Getting Started

## Alternative installations



```
dnf install hypershell (EPEL 10)
yum install hypershell
```

```
apt install hypershell (coming soon)
snap install hypershell (coming soon)
```

```
apptainer build ...
```

```
docker build ... (ghcr.io/hypershell/hypershell) **
```

# *Getting Started*

## Alternative installations

Reach out to me for details!

Ask your HPC admins to add HyperShell as a module!

LMOD module available

```
$ module avail hypershell
```

```
----- Core Applications -----  
hypershell/2.7.0 (D)
```

Where:

D: Default Module

# 3

## Basics

The 'hello world' of workflows. Basic end-to-end.

# Basics

[https://github.com/NERSC/htc\\_training/tree/main/hypershell/00\\_first](https://github.com/NERSC/htc_training/tree/main/hypershell/00_first)

## Hello World

- Let's run the simplest possible example to get started.
- The `hsx` program is short-hand for `hs cluster` and `-t` applies a command template to incoming arguments – not unlike *GNU Parallel*.

Simple 'hello world' example

```
$ seq 4 | hsx -t 'echo {}'  
WARNING [hypershell.server] No database configured - automatically disabled  
1  
2  
3  
4
```



Use '`hs --help`' or '`man hs`' for help and detailed usage information. How can we disable this warning?

# Basics

[https://github.com/NERSC/htc\\_training/tree/main/hypershell/01\\_basic](https://github.com/NERSC/htc_training/tree/main/hypershell/01_basic)

## Parallel execution of tasks

- The `-N` flag defines the number of parallel works (task executors).
- Later, we'll see this is the number of workers *per-client*.

### Parallel execution of tasks

```
$ seq 12 | hsx -N4 -t 'sleep 1 && echo {}' --no-db --no-confirm  
1  
2  
3  
4  
...
```



Notice while this runs that the output comes back in bursts of 4.

# Basics

[https://github.com/NERSC/htc\\_training/tree/main/hypershell/02\\_basic\\_templates](https://github.com/NERSC/htc_training/tree/main/hypershell/02_basic_templates)

## Template patterns

- Rich template pattern syntax for quickly mapping inputs to commands.
- Inspired by *GNU Parallel*.
- See [documentation](#) for reference.

Template patterns for mapping inputs to outputs

```
$ find in/ -type f | hsx -N2 -t 'grep NEEDLE {} > out/{/-}.out' --no-db --no-confirm
```



Select out arguments from input lines using `{[2]}` or `{[-3:]}` slicing like in Python.

# Basics

[https://github.com/NERSC/htc\\_training/tree/main/hypershell/03\\_basic\\_failures](https://github.com/NERSC/htc_training/tree/main/hypershell/03_basic_failures)

## Capturing failed tasks

- Use **-f/--failures** to redirect inputs that result in a non-zero exit status.
- We're simulating failures using the Unix **true** and **false** programs.
- Later, we'll see how to use **--max-retries** with the database enabled.

### Capturing failed task inputs

```
$ hsx task.in -N4 --no-db --no-confirm -f task.failed
...
WARNING [hypershell.server] Non-zero exit status (1) for task (4806ada6-7b47-
47e9-b78a-f6aa042a830b)
```



We can cycle the failed task file back in as input for a crude retry mechanism – not unlike *ParaFly*.

# Basics

[https://github.com/NERSC/htc\\_training/tree/main/hypershell/04\\_basic\\_slurm](https://github.com/NERSC/htc_training/tree/main/hypershell/04_basic_slurm)

## Slurm job script

- Slurm is the most common scheduler at HPC centers.
- This is a single-node job running single-core tasks.

```
Text Editor

#!/bin/bash
#SBATCH -A mylab -p cpu -q normal
#SBATCH -c42 -t 1-00:00:00

module load hypershell

hsx task.in -N42 -f task.failed --no-db --no-confirm
```



# Basics

[https://github.com/NERSC/htc\\_training/tree/main/hypershell/05\\_basic\\_capture\\_output](https://github.com/NERSC/htc_training/tree/main/hypershell/05_basic_capture_output)

## Capturing output

- We'll learn more about configuration soon.
- Use **HYPERHELL\_SITE** to control where outputs are directed with **--capture**

```
Text Editor

#!/bin/bash
#SBATCH -A mylab -p cpu -q normal
#SBATCH -c42 -t 1-00:00:00

module load hypershell

export HYPERHELL_SITE=$(pwd)
hsx task.in -N42 -f task.failed --no-db --no-confirm --capture
```

# Basics

[https://github.com/NERSC/htc\\_training/tree/main/hypershell/06\\_basic\\_multinode](https://github.com/NERSC/htc_training/tree/main/hypershell/06_basic_multinode)

## Distributed computing (many node)

- Let's do the same thing again but with many tasks (100k) across many nodes.
- The **--launcher=srun** brings up the **hs client** on each node.

```
Text Editor

#!/bin/bash
#SBATCH -A mylab -p cpu -q normal
#SBATCH -N8 -c192 --exclusive -t 1-00:00:00

module load hypershell

export HYPERHELL_SITE=$(pwd) HYPERHELL_LOGGING_LEVEL=DEBUG
hsx task.in -N192 -b192 --launcher=srun -f task.failed \
    --no-db --no-confirm 1>task.out 2>task.log
```

# Basics

[https://github.com/NERSC/htc\\_training/tree/main/hypershell/07\\_extreme\\_scale](https://github.com/NERSC/htc_training/tree/main/hypershell/07_extreme_scale)

## Extreme scale computing

- Use **-w / --bundlewait** to govern task bundle synchronization (avoid under-filled bundles).
- The **--delay-start** option staggers the launch across a large cluster.

```
Text Editor

#!/bin/bash
#SBATCH -A mylab -p cpu -q normal
#SBATCH -N1000 -c192 --exclusive -t 1-00:00:00

module load hypershell

export HYPERSHELL_SITE=$(pwd) HYPERSHELL_LOGGING_LEVEL=DEBUG
hsx task.in -N192 -b192 -w60 --launcher=srun -f task.failed \
    --no-db --no-confirm --delay-start=-30 1>task.out 2>task.log
```

# 4

## Configuration

Configuration. Submitting to a database.

# *Configuration*

## Setting parameters globally

- Define parameters globally that persist between invocations.
- Set from the command-line or edit the file directly.

```
$ hs config --help
$ hs config set logging.level debug --user

$ hs config which logging.level
$ hs config edit --user
```

# Configuration

## Local SQLite database

- Include a database path and HyperShell will persist tasks.
- Be careful which filesystem this points to (e.g., /tmp or /home is best).

```
Text Editor

# File automatically created on 2025-05-28 19:05:51.383195
# Settings here are merged automatically with defaults and environment variables

[logging]
level = "debug"

[database]
file = "/home/glentner/.hypershell/lib/main.db"
```

# Configuration

[https://github.com/NERSC/htc\\_training/tree/main/hypershell/08\\_database\\_intro](https://github.com/NERSC/htc_training/tree/main/hypershell/08_database_intro)

## Submitting tasks to the database

- The server/cluster need not be running to add tasks.
- Use **hs submit** with either single commands or collections from a file.

### Submit tasks

```
$ hs submit echo "hello world"
DEBUG [hypershell.submit] Submitted single task (explicit)
INFO [hypershell.submit] Submitted task (ce8893cd-c587-4f44-9499-7c679bd2b437)

$ hs submit <(seq 100) --template "echo {}"
DEBUG [hypershell.submit] Submitted from /proc/self/fd/11 (implicit - not executable)
...
INFO [hypershell.submit] Submitted 100 tasks
```



What command-line option can we use to make batch file submission “explicit”?

# Restarts

Submit tasks now - restart the cluster later

- It is often better to have your workflow “restart” even from the beginning.
- If there are no tasks remaining the program will simply shut down.

```
Restart cluster where it left off

$ hsx --restart ...
  DEBUG [hypershell.server] Started
...
  WARNING [hypershell.server] Database exists (101 previous tasks)
  INFO [hypershell.server] Found 101 unfinished task(s)
  INFO [hypershell.server] Reverted 0 previously interrupted task(s)
  DEBUG [hypershell.server] Scheduled 1 tasks
  DEBUG [hypershell.server] Scheduled task (ce8893cd-c587-4f44-9499-7c679bd2b437)
...
```



# 5

## GPU Multiplexing

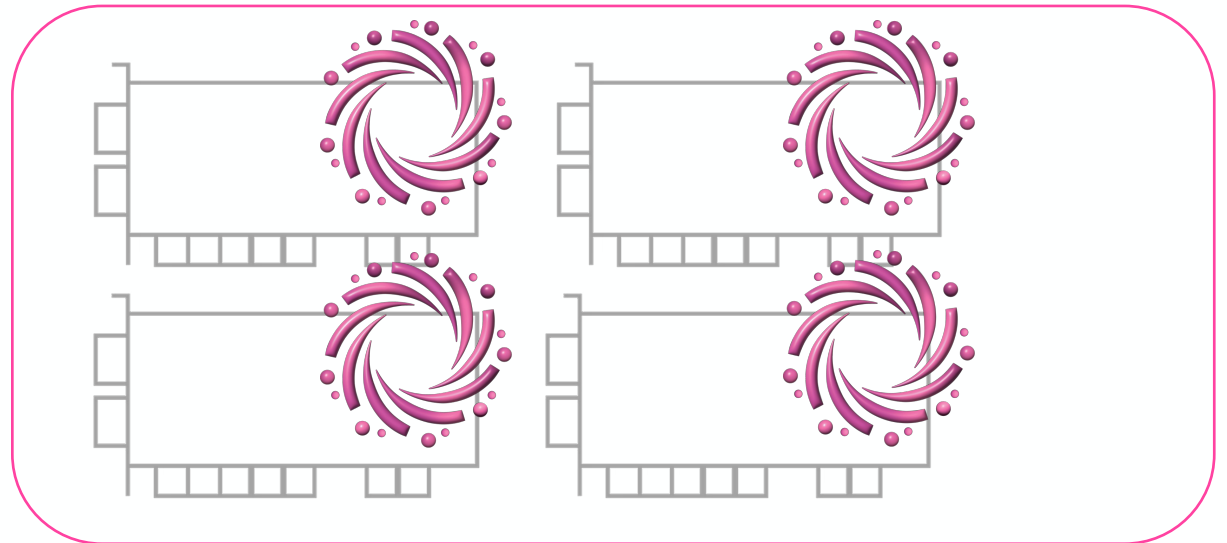
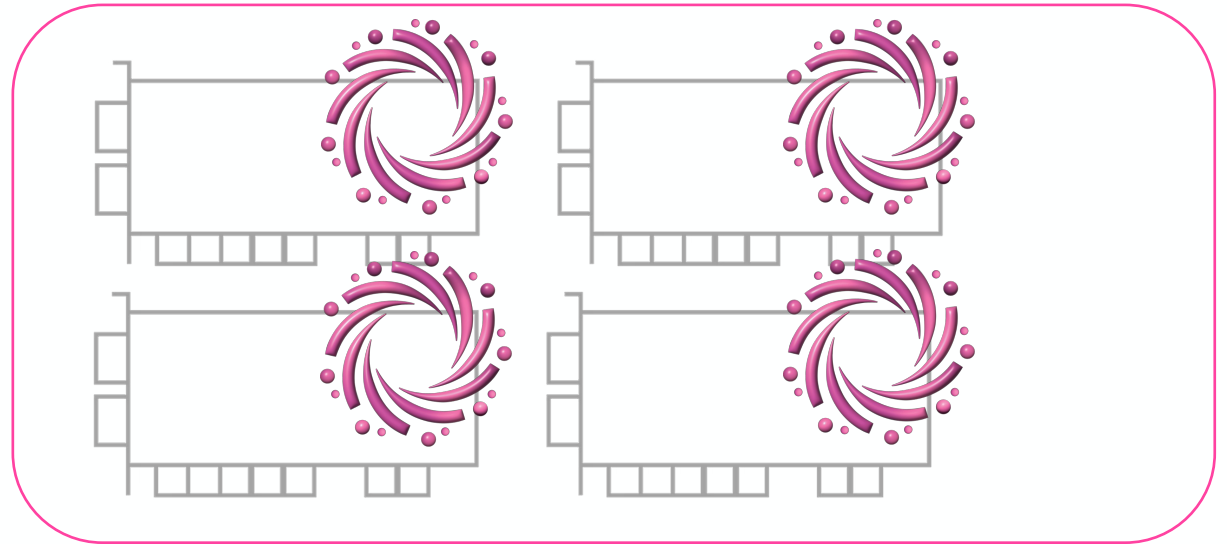
Specific use-case in demand right now.

# ***GPU Multiplexing***

## Packing small tasks

- GPUs are expensive!
- Make better use of large GPUs.
- Accelerate smaller tasks in parallel with one client per GPU.
- Or deploy persistent autoscaling cluster for throughput training, more on this next!

Node 1



Node 2

# *GPU Multiplexing*

## Packing small tasks correctly

- Large GPUs are expensive – researchers still operate small models/tasks.
- One Slurm task per GPU launches one **hs client** per GPU.

```
Text Editor

#!/bin/bash
#SBATCH -A mylab -p gpu -q normal
#SBATCH --gres=gpu:2 --tasks-per-gpu=1 -t 1-00:00:00

module load hypershell

export HYPERHELL_SITE=$(pwd) HYPERHELL_LOGGING_LEVEL=DEBUG
hsx task.in -N4 --launcher=srun -f task.failed \
    --no-db --no-confirm 1>task.out 2>task.log
```

# 6

## Server vs Client

Running the server separate from clients.  
Submitting tasks.

# Server vs Client

## One server many clients

- Keep the server alive elsewhere and scale out clients as needed.
- Protect using secure auth key! (automatic with cluster)
- Bind to 0.0.0.0 to allow remote connections (default: localhost)

### Run server

```
$ hs server --forever -k mykey123 --bind 0.0.0.0 --max-retries 2
DEBUG [hypershell.server] Started
DEBUG [hypershell.server] Started (scheduler)
INFO [hypershell.server] Scheduler will run forever
...
WARNING [hypershell.server] Database exists (101 previous tasks)
WARNING [hypershell.server] All tasks completed - did you mean to use the same database?
DEBUG [hypershell.server] Registered client (login07.gautschi.rcac.purdue.edu: ...)
DEBUG [hypershell.server] Checking clients (1 connected)
...
```

# Server vs Client

## Clients connect and process tasks

- Clients register themselves with the server and process tasks locally.
- Clients can be evicted if they miss too many heartbeats. (tasks are re-scheduled)
- Set a timeout to have them shutdown if there are no tasks. (default: never disconnect)

Run server

```
$ hs client -N16 --host login07 -k mykey123 --timeout 60 --capture
DEBUG [hypershell.client] Started (16 executor)
DEBUG [hypershell.client] Started (scheduler: no timeout)
DEBUG [hypershell.client] Started (collector)
DEBUG [hypershell.client] Started (heartbeat)
DEBUG [hypershell.client] Started (executor-1)
...
```

# 7

## Task Management

Task search and updates.

Cancelling, reverting, and deleting tasks.

User-defined tags (command, inline, group).

# Task Management

## Search database for task history

- Use **hs list** to search task history and remaining tasks.
- Many output formats (normal, list, table, csv, json).
- Output file paths stored in database for retrieval.

```
List task history

$ hs list -l1
---
      id: ce8893cd-c587-4f44-9499-7c679bd2b437
     args: echo "hello world"
...
  started: 2025-05-29 21:32:02.041907 (waited: 0:05:22)
 completed: 2025-05-29 21:32:02.045961 (duration: null)
submit_host: login07.gautschi.rcac.purdue.edu (f5a671b9-d7af-4b9c-aa5c-44e41b332a96)
...
     tags: part:0
```



# *Task Management*

## Search database for task history

- Find most recent completed tasks, all remaining tasks, by client host, etc ...
- User defined tags allow for custom attributes for single or groups of tasks.

Count of remaining tasks within group belonging to tag

```
$ hs list --count --remaining --with-tag site:b  
13
```



Including only the tag “key” will return any task with that key regardless of value.

# *Task Management*

## User-defined tags

- Apply tags at submit time from the command-line.
- All tasks inherit these tags.

Submit group of tasks with shared tag

```
$ seq 1000 | hs submit -t group:1 demo  
...
```

# *Task Management*

## Inline tag assignments

- Include tag assignments with special comment syntax within submission files.
- Combine these two approaches together in one invocation.

Input task file with inline tag comments

```
echo 1  # HYPERHELL: data:1
echo 2  # HYPERHELL: data:2
echo 3  # HYPERHELL: data:3
echo 4  # HYPERHELL: data:4
echo 5  # HYPERHELL: data:5
echo 6  # HYPERHELL: data:6

...
```

# *Task Management*

## Inline tag assignments

- Stand-alone comments apply tags to all tasks. (until overridden)
- Combine these three approaches together in one invocation.

Input task file with inline tag comments

```
# HYPERHELL: case:1
echo AA
echo BB

# HYPERHELL: case:2
echo CC
echo DD

...
```

# 8

## Autoscaling

Elastic scaling the cluster - even to zero.

# Autoscaling

## Dynamically launch clients

- Based on “task pressure” (dimensionless quantity)
- The launcher is used as a prefix to provision a client.
- Policy based with initial / min / max sized cluster.
- Set a client timeout and send SIGUSR1 ahead of job timeout to drain the worker.

Run autoscaling cluster

```
$ hsx -N192 -P5 -I1 -X0 -Y12 -T60 --capture --autoscaling=dynamic \  
    --launcher="srun -Q -A mylab -p cpu -t60 -c192 --signal=SIGUSR1@600"  
...
```

# 9

## Automated Clusters

User facility (service) deployments.  
Automatic scaling of pilot jobs.  
Database partitioning.

# Automated Clusters

## Service account pilot jobs

- Run dedicated server with *PostgreSQL* off-cluster (e.g., Kubernetes).
- Submit **xcore-pilot** job to run clients, as necessary.

### Pilot job

```
#!/bin/bash
#SBATCH -A xcore -p cpu -q normal -J xcore-pilot
#SBATCH -N1 -n1 -c48 -t 01:00:00
#SBATCH -o /dev/null
#SBATCH --signal B:USR1@300

exec ~/bin/xcore-hs client -N48 -T60 --capture \
    --host hs-server.xcore.university.edu -k `cat ~/.hypershell/key` \
    2>>$SCRATCH/xcore-pilot/log/hypershell-`date +%Y%m%d`.log
```



# Automated Clusters

Submit jobs only when necessary

- Cron (or similar) execution of **xcore-autoscale** script.
- The **xcore-hs** script is just a wrapper around **hs** itself.

Pilot job

```
#!/bin/bash

JOBS=$(/usr/bin/sacct -n -X -u xcore -s R,PD --name xcore-pilot | wc -l)
TASKS=$(~/bin/xcore-hs list --count --remaining)

if [ $JOBS -gt 8 ] || [ $(( (TASKS / 48) - JOBS )) -lt 1 ]
then
    echo "no scale"
else
    sbatch ~/bin/xcore-pilot
fi
```

*So much more we could dive into,  
but this covers the gist of it.*

# *End*

Please fill out the survey!

Questions?